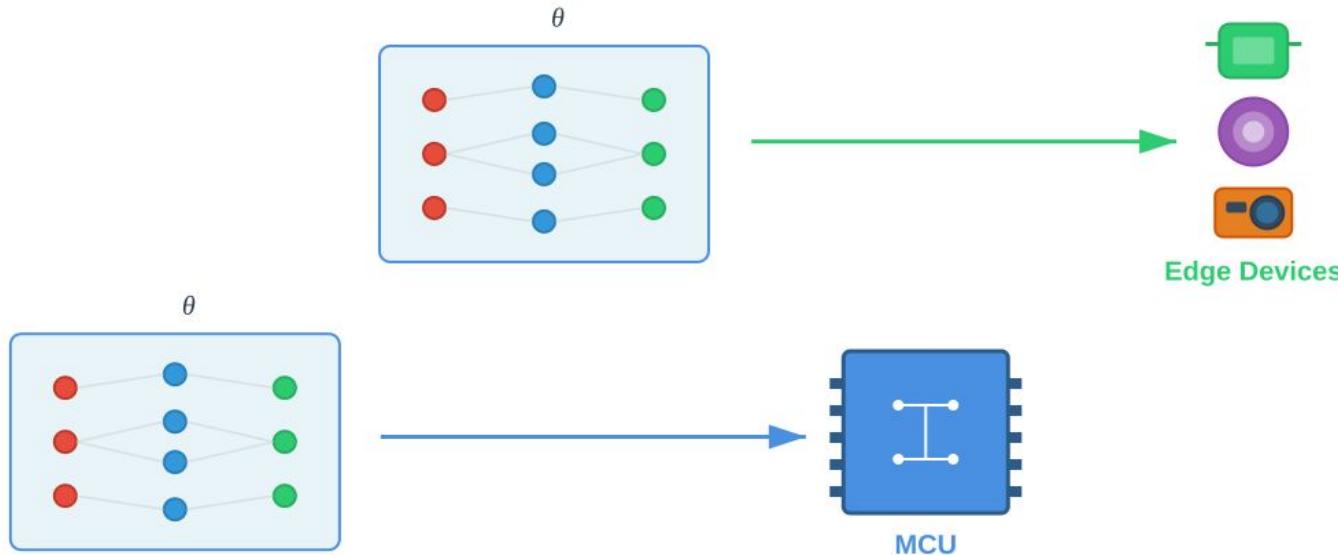


Edge AI & TinyML



Bringing Intelligence to the Edge

Presented by
Dr. Mohammed Alnemari

Dr. Mohammed Alnemari

- ASSISTANT PROFESSOR IN COMPUTER ENGINEERING SPECIALIZING IN ARTIFICIAL INTELLIGENCE IN EMBEDDED SYSTEMS
- ASSOCIATE RESEARCHER AT THE ARTIFICIAL INTELLIGENCE AND SENSING TECHNOLOGIES RESEARCH CENTER.
- AI CONSULTING AND ADVANCED TECHNOLOGIES

Educational Background:

1. Bachelor's in Computer Engineering from Taif University (Saudi Arabia)
2. Bridge Program at the University of Pennsylvania (USA)
3. Master's degree from the University of California, Irvine (USA)
4. Visiting Master's student at National Taiwan University (Taiwan)
5. Visiting PhD student at Tokyo Institute of Technology (Japan)
6. PhD from the University of California, Irvine (USA)
7. Visiting Researcher at the University of Southern California (USA)

RESEARCH INTERESTS: EDGE AI, TinyML , NEURAL NETWORK COMPRESSION, EFFICIENT NEURAL NETWORKS, EMBEDDED MACHINE LEARNING



TABLE OF CONTENT

2 MOTIVATION & CONTEXT

3 CORE CONCEPTS

4 TECHNICAL FOUNDATIONS

5 HANDS-ON UNDERSTANDING

6 TOOLS AND ECOSYSTEM

7 BROADER IMPLICATIONS

8 LIVE DEMONSTRATION

9 CONCLUSION

WHY EDGE AI?

The Data Explosion Challenge



Three Critical Limitations of Cloud Computing



1. Latency

- Network transmission delays
- Unusable for real-time
- $T = T_{\text{transmit}} + T_{\text{compute}}$



2. Privacy

- Sensitive data exposure
- Regulatory compliance
- Data sovereignty concerns

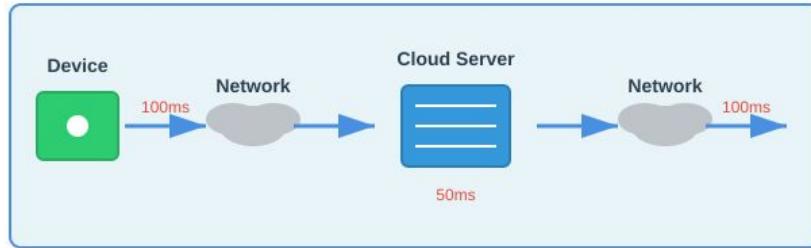


3. Bandwidth

- Network congestion
- High operational costs
- Limited remote connectivity

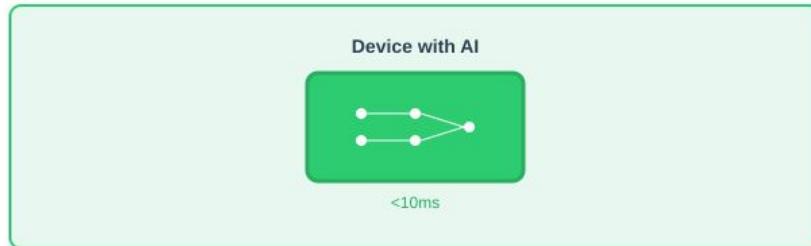
FROM CLOUD AI → EDGE AI

Cloud Inference Architecture



Total Latency: ~250ms

Edge Inference Architecture



Total Latency: <10ms

Mathematical Cost

$$T = T_{transmit} + T_{compute}$$

Where:

$$T_{transmit} = \text{Data_size} / \text{Bandwidth}$$

$$T_{compute} = \text{Operations} / \text{FLOPS}$$

Critical for Control Loops

Latency > 100ms can be catastrophic:

- Autonomous vehicles
- Industrial robotics
- Medical devices
- Energy grids

EDGE AI APPLICATIONS



Smart Home

- Voice assistants
- Security cameras
- Smart thermostats



Wearables

- Health monitoring
- Fitness tracking
- Fall detection



Agriculture

- Crop monitoring
- Soil analysis
- Irrigation control



Industrial IoT

- Predictive maintenance
- Quality control
- Process optimization



Autonomous Systems

- Self-driving vehicles
- Drones
- Navigation systems



Energy Systems

- Smart grids
- Solar monitoring
- Battery management

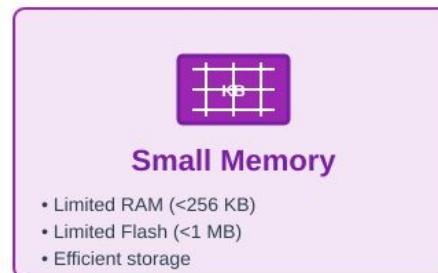
WHAT IS EDGE AI?

CORE CONCEPTS

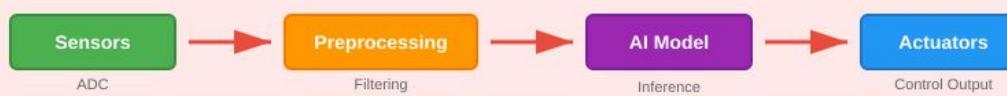
Definition

Edge AI refers to artificial intelligence algorithms running locally on embedded or IoT devices at the "edge" of the network, rather than in centralized cloud servers.

Key Requirements



Edge AI Architecture



WHAT IS TINYML?

CORE CONCEPTS

Definition

TinyML is the deployment of machine learning models on microcontrollers (MCUs) and ultra-low-power embedded systems.

Typical MCU Specifications

| Component | Specification |
|-------------------|-----------------------------------|
| Processor | 32-bit ARM Cortex-M4 |
| Clock Speed | 64-80 MHz |
| Flash Memory | 512 KB - 1 MB |
| RAM | 128-256 KB |
| Power Consumption | <1 mW (sleep), 50-100 mW (active) |
| Cost | \$2-\$10 |

TinyML Capabilities: Audio keyword spotting • Gesture recognition • Anomaly detection • Image classification • Sensor fusion

Community & Standards: IEEE TinyML community • MLPerf Tiny benchmark • TinyML Foundation

TINYML VS EDGE AI VS CLOUD AI

Comparative Analysis

| Feature | Cloud AI | Edge AI | TinyML |
|-----------------|---------------------|----------------------|---------------------|
| Hardware | GPU/TPU Servers | Edge TPU, Mobile GPU | Microcontrollers |
| Memory | GBs to TBs | GBs | KBs to MBs |
| Power | kW | Watts | mW to μ W |
| Latency | 100-500ms | 10-50ms | <10ms |
| Connectivity | Required | Optional | None |
| Privacy | Low (data sent) | Medium | High (local) |
| Model Size | Billions of params | Millions of params | Thousands of params |
| Cost per device | Subscription | \$50-\$200 | \$2-\$10 |
| Update Ability | Easy | Moderate | Challenging |
| Applications | Complex NLP, Vision | Autonomous vehicles | Sensors, wearables |

CHALLENGES IN BRINGING AI TO THE EDGE

Memory

Power

Learning

TECHNICAL FOUNDATIONS

1. Limited Compute and Memory Resources

Challenge:

- Modern DNNs: millions to billions of parameters
- MCU Flash: typically <1 MB
- MCU RAM: typically <256 KB

Mathematical Constraint:

$$\begin{aligned} \text{Model_size} &\leq \text{Flash_available} \\ \text{Activation_size} &\leq \text{RAM_available} \end{aligned}$$

Example:

- MobileNetV2: ~14 MB (uncompressed)
- Target MCU: 1 MB Flash
- Compression ratio needed: 14x

2. Power Consumption Constraints

Energy Budget:

- Coin cell battery: ~600 mAh at 3V = 6.5 kJ
- Target lifetime: 1 year = 31,536,000 seconds
- Available power: ~200 µW continuous

Energy per Inference:

$$\begin{aligned} E_{\text{inference}} &= P_{\text{compute}} \times t_{\text{inference}} \\ \text{For 100 inferences/day:} \\ E_{\text{daily}} &= 100 \times E_{\text{inference}} \\ E_{\text{daily}} &\leq E_{\text{budget_daily}} \end{aligned}$$

Typical Power Budget:

Active: 10-100 mW | Sleep: <10 µW | Average: <1 mW

3. On-Device Learning Limitations

Challenges:

- Backpropagation requires storing activations
- Gradient computation memory-intensive
- Limited training data on device
- Energy cost of training 1000x higher than inference

Solutions:

- ✓ Transfer learning
(fine-tune pre-trained models)
- ✓ Online learning
(incremental updates)
- ✓ Federated learning
(collaborative without data sharing)

Challenge Interconnections:



OBJECTIVES TO ACHIEVE

OPTIMIZATION GOALS



Minimize number of the weights and activations in the networks.



Minimize the time of the inference.



Minimize the memory complexity



Minimize the FLOP operation (floating point operation per second).



Minimize the energy requirements.

THE ECOSYSTEM OF THE SOLUTION

TOOLS AND ECOSYSTEM



Compression Algorithms

1. Quantization

Reduce precision (FP32 → INT8) to minimize model size and accelerate inference.

2. Pruning

Remove redundant weights to create sparse networks with minimal accuracy loss.

3. Knowledge Distillation

Train small student model from large teacher to transfer learned knowledge.

4. Tensor Decomposition

Factorize weight tensors into low-rank components for parameter reduction.

System-Level Tools

TVM

Deep Learning Compiler

MLIR

Multi-Level IR

ONNX Runtime

Inference Optimizer

TFLite Micro

MCU Inference

Custom Hardware

FPGA

FPGA

ASIC

Note: Frameworks like PyTorch/TensorFlow provide high-level APIs, while system tools optimize for hardware deployment.

TINYML HARDWARE PLATFORMS

HANDS-ON UNDERSTANDING

Popular Development Boards



Arduino Nano 33 BLE Sense

- Cortex-M4 @ 64 MHz
- 1 MB Flash, 256 KB RAM
- IMU, Mic, Temp, Humidity



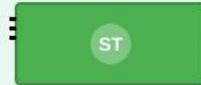
Raspberry Pi Pico

- Dual Cortex-M0+ @ 133 MHz
- 2 MB Flash, 264 KB RAM
- Low cost (\$4)



ESP32

- Dual-core Tensilica @ 240 MHz
- 4 MB Flash, 520 KB RAM
- WiFi + Bluetooth



STM32 Nucleo

- Cortex-M4 @ 80 MHz
- 1 MB Flash, 128 KB RAM
- Arduino-compatible



Sony Spresense

- 6-core Cortex-M4F @ 156 MHz
- 8 MB Flash, 1.5 MB RAM
- GPS, Audio processing



Nvidia Jetson Nano

- Quad-core ARM A57 @ 1.43 GHz
- 128-core Maxwell GPU
- Edge AI (not TinyML)

Selection Criteria:

Power budget • Memory requirements • Sensor integration • Connectivity • Cost • Development ecosystem

NEURAL NETWORK COMPRESSION

TECHNICAL FOUNDATIONS

Challenge: Modern neural networks are too large for edge devices.

Solution: Compress models while maintaining accuracy.

Five Main Compression Techniques



1. Quantization

- Reduce precision: FP32 → INT8
- 4x smaller model, 4x faster
- Minimal accuracy loss (<1%)



2. Pruning

- Remove unimportant connections
- Structured vs. unstructured
- Up to 90% sparsity possible

3. Knowledge Distillation



- Train small from large model
- Transfer "soft" predictions
- Better than from scratch

5. Tensor Decomposition



- Low-rank factorization
- Tucker, CP decomposition
- Reduce parameter count

4. Efficient Architectures



- MobileNet, EfficientNet
- Depthwise separable conv.
- Efficient from start

QUANTIZATION DEEP DIVE

TECHNICAL FOUNDATIONS

What is Quantization?

Mapping high-precision values (FP32) to low-precision values (INT8, INT4)

Quantization Formula

$$q = \text{round}(r / S) + Z$$

S = scale, Z = zero-point, r = real value, q = quantized

Precision Comparison

FP32 (32-bit Floating Point)



Benefits & Trade-offs

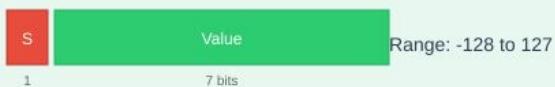
✓ Benefits

- 4x reduction in model size
- 2-4x faster inference
- Lower memory bandwidth

⚠ Trade-offs

- Accuracy loss (typically <1%)
- Requires calibration data
- Hardware support needed

INT8 (8-bit Integer)



Types:

Post-Training Quantization

Quantization-Aware Training

Dynamic Range

Popular Frameworks: TensorFlow Lite • PyTorch Mobile • ONNX Runtime • TFLite Micro

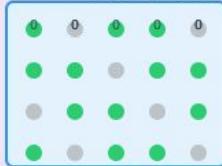
PRUNING DEEP DIVE

TECHNICAL FOUNDATIONS

Pruning removes redundant or less important weights/neurons from neural networks to create sparse models while maintaining accuracy.

Two Main Types of Pruning

Unstructured Pruning



- Remove individual weights
- Irregular sparsity pattern
- Requires sparse matrix libraries

Structured Pruning



- Remove entire filters/channels
- Regular sparsity pattern
- Hardware-friendly (no special libs)

Typical Pruning Process

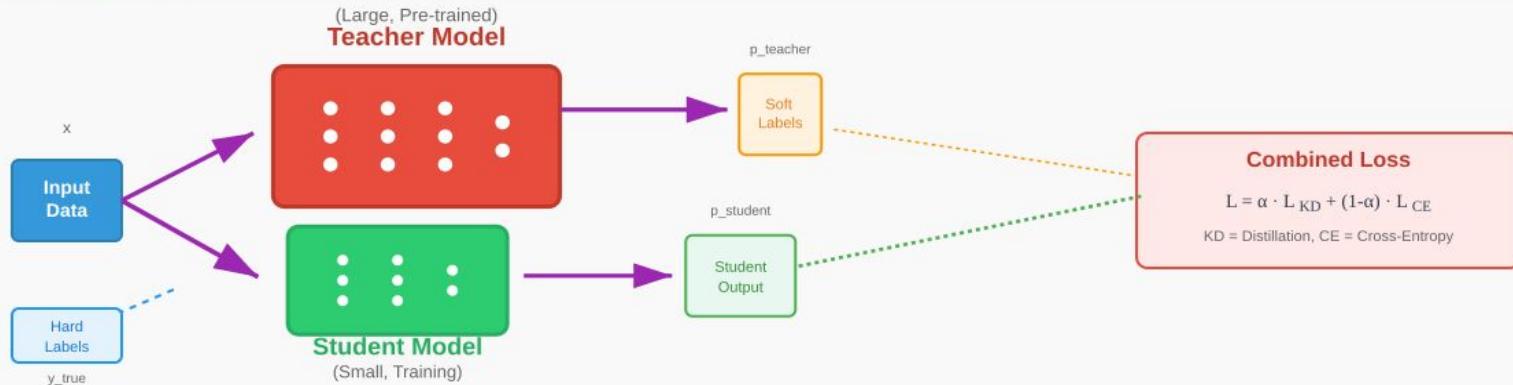
1. Train Model
2. Identify Weights
3. Prune Weights
4. Fine-tune
5. Repeat/Deploy

KNOWLEDGE DISTILLATION

TECHNICAL FOUNDATIONS

Knowledge Distillation transfers knowledge from a large "teacher" model to a smaller "student" model, achieving better performance than training from scratch.

Architecture



Temperature Softmax

$$p_i = \exp(z_i/T) / \sum \exp(z_j/T)$$

T = Temperature (higher = softer probabilities)

Typical: T = 3-5 during training, T = 1 at inference

✓ Key Advantages

- Student learns from teacher's soft predictions
- Better generalization than training alone
- Can compress 10-100x with minimal accuracy loss

TENSOR DECOMPOSITION

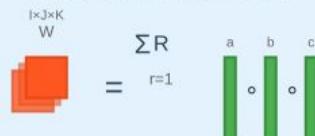
TECHNICAL FOUNDATIONS

Tensor Decomposition decomposes high-dimensional weight tensors into products of smaller tensors, reducing parameters and computation.

Main Decomposition Methods

CP Decomposition

(CANDECOMP/PARAFAC)

$$\underset{I \times J \times K}{W} = \sum_{r=1}^R a \circ b \circ c$$
A diagram illustrating CP decomposition. On the left, a 3D tensor W is shown as a stack of three red squares. To its right is the equation $\underset{I \times J \times K}{W} = \sum_{r=1}^R a \circ b \circ c$. To the right of the equation are three vertical green bars labeled 'a', 'b', and 'c' with small circles between them, representing the factor matrices.

- Sum of rank-1 tensors
- R = rank (hyperparameter)
- Compression: IJK → R(I+J+K)

Tucker Decomposition

$$\underset{I \times J \times K}{W} = \underset{R_1 \times R_2 \times R_3}{\text{Core}} \times \underset{A \times B \times C}{\text{A, B, C}}$$
A diagram illustrating Tucker decomposition. On the left, a 3D tensor W is shown as a stack of three red squares. To its right is the equation $\underset{I \times J \times K}{W} = \underset{R_1 \times R_2 \times R_3}{\text{Core}} \times \underset{A \times B \times C}{\text{A, B, C}}$. To the right of the equation are three vertical green bars labeled 'A', 'B', and 'C' with small circles between them, representing the factor matrices. The core tensor is shown as a purple cube.

- Core tensor + factor matrices
- More flexible than CP
- Compression: IJK → $R_1R_2R_3 + IR_1 + JR_2 + KR_3$

SVD (Matrix Case)

Singular Value Decomposition

$$\underset{m \times n}{W} = \underset{m \times r}{U} \underset{\Sigma \text{ r} \times r}{\Sigma} \underset{r \times n}{V^T}$$
A diagram illustrating SVD decomposition. On the left, a matrix W is shown as a red rectangle. To its right is the equation $\underset{m \times n}{W} = \underset{m \times r}{U} \underset{\Sigma \text{ r} \times r}{\Sigma} \underset{r \times n}{V^T}$. To the right of the equation are two rectangles: a blue one labeled 'U' and a purple one labeled 'V^T'. Between them is a yellow square labeled 'Sigma' with diagonal dots.

- Low-rank approximation
- Keep top r singular values
- Compression: mn → r(m+n+1)

Practical Example: Conv Layer Decomposition

Conv2D: $[C_{\text{out}} \times C_{\text{in}} \times K \times K]$ tensor → Decomposed into smaller filters

Example: $256 \times 256 \times 3 \times 3$ (589K params) → $256 \times 64 + 64 \times 256 \times 3 \times 3$ (163K params) ≈ 72% reduction

MODEL DEPLOYMENT FRAMEWORKS

TINYML ECOSYSTEM

Popular Deployment Frameworks



TensorFlow Lite Micro

Features:

- C++11 runtime, no OS required
- Memory optimized (<20 KB)
- 90+ operators supported
- Op-level profiling built-in

Best for:

Production deployments, ARM Cortex-M



Edge Impulse Studio

Features:

- End-to-end ML platform
- Data collection & labeling
- AutoML capabilities
- Deploy to 100+ boards

Best for:

Rapid prototyping, audio/vision tasks



PyTorch Mobile / ONNX

Features:

- PyTorch ecosystem integration
- ONNX interoperability
- Mobile & edge optimized
- Growing MCU support

Best for:

Research projects, cross-platform

Typical Deployment Pipeline



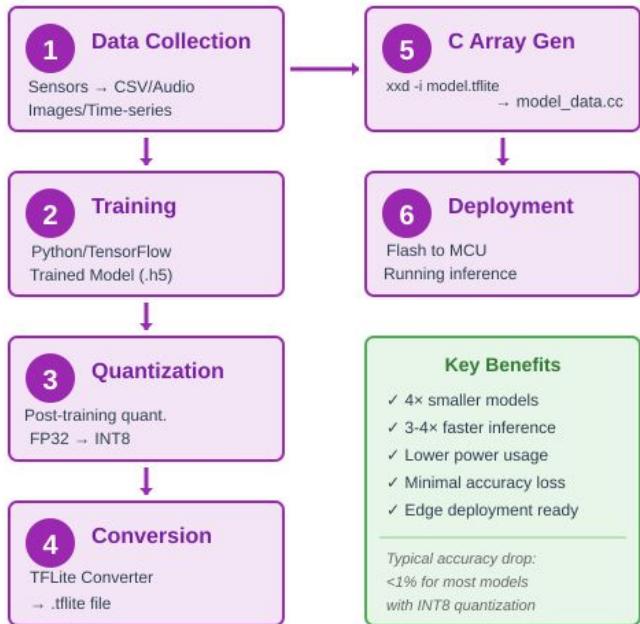
Key Considerations:

Ease of use • Hardware support • Model format • Community • Documentation • Performance optimization

FROM MODEL TO MICROCONTROLLER

HANDS-ON UNDERSTANDING

Complete Workflow



Code Example: Conversion to TFLite

```
# Convert to TFLite
converter = tf.lite.TFLiteConverter.from_keras_model(model)
converter.optimizations = [tf.lite.Optimize.DEFAULT]
tflite_model = converter.convert()

# Save model
with open('model.tflite', 'wb') as f:
    f.write(tflite_model)

# Generate C array
$ xxd -i model.tflite > model_data.cc

# Result: C array ready for MCU
unsigned char model_tflite[] = {
    0x1c, 0x00, 0x00, 0x00, 0x54, 0x46, ...};

unsigned int model_tflite_len = 10240;
```

- Key Benefits**
- ✓ 4x smaller models
 - ✓ 3-4x faster inference
 - ✓ Lower power usage
 - ✓ Minimal accuracy loss
 - ✓ Edge deployment ready

*Typical accuracy drop:
<1% for most models
with INT8 quantization*

Model size reduced: FP32 (40KB) → INT8 (10KB) → 4x compression

Inference speed: ~8ms on Cortex-M4 @ 64MHz

Memory footprint: 12KB RAM, 10KB Flash

Note: Always profile your model on target hardware to verify performance meets requirements

CASE STUDY: GESTURE RECOGNITION

REAL-WORLD APPLICATION

Problem Statement

Objective: Recognize hand gestures using accelerometer data on MCU

Gestures:

- Wave (side to side motion)
- Fist (punch motion)
- Open hand (spread fingers)
- Idle (no motion)

Dataset

Sensor: 3-axis accelerometer (LSM9DS1)

Sampling rate: 119 Hz

Window size: 1 second (119 samples)

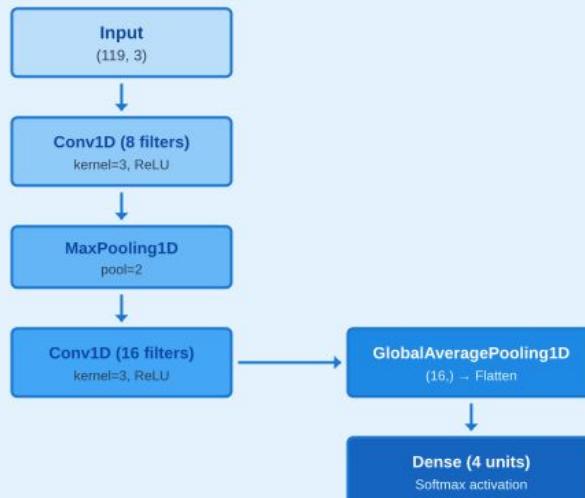
Features: $ax, ay, az \times 119 = 357$ inputs

Samples: 1000 per gesture class

Sample data visualization:



Model Architecture



Model Statistics

Total parameters: ~2,500

FP32 size: 10 KB

INT8 size: 2.5 KB

Trainable layers: 4

Input shape: (119, 3)

Output shape: (4,)

✓ **Lightweight**

✓ **Fast inference**

✓ **MCU-friendly**

Performance Metrics

| Metric | Value |
|--------------------|-------|
| Accuracy | 94.5% |
| Inference time | 8 ms |
| Memory (RAM) | 12 KB |
| Model size (Flash) | 10 KB |
| Power consumption | 45 mW |

Confusion Matrix (simplified)

| | Wave | Fist | Open | Idle |
|------|------|------|------|------|
| Wave | 95 | 2 | 2 | 1 |
| Fist | 3 | 93 | 3 | 1 |
| Open | 1 | 4 | 94 | 1 |
| Idle | 0 | 1 | 2 | 97 |

CODE SNIPPET: PYTHON → C++

C++ Inference Code

FROM TRAINING TO DEPLOYMENT

Python Training Code

```
import tensorflow as tf

# Define model
model = tf.keras.Sequential([
    tf.keras.layers.Conv1D(8, 3, activation='relu', input_shape=(119, 3)),
    tf.keras.layers.MaxPooling1D(2),
    tf.keras.layers.Conv1D(16, 3, activation='relu'),
    tf.keras.layers.GlobalAveragePooling1D(),
    tf.keras.layers.Dense(4, activation='softmax')])

# Compile
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])

# Train
model.fit(X_train, y_train,
          epochs=50, validation_split=0.2)

# Convert to TFLite
converter = tf.lite.TFLiteConverter.from_keras_model(model)
```



```
// Include TFLite Micro
#include "tensorflow/lite/micro/
micro_interpreter.h"
#include "model_data.h"

// Allocate memory
constexpr int kTensorArenaSize =
20 * 1024;
uint8_t tensor_arena[kTensorArenaSize];

// Setup interpreter
tflite::MicroInterpreter interpreter(
    model, resolver,
    tensor_arena, kTensorArenaSize);

// Allocate tensors
interpreter.AllocateTensors();

// Get input tensor
TfLiteTensor* input =
interpreter.input(0);

// Copy sensor data
for (int i = 0; i < 357; i++) {
    input->data.int8[i] =
    quantize(sensor_data[i]);
}
```

```
// Run inference
interpreter.Invoke();

// Get output and process result
TfLiteTensor* output = interpreter.output(0);
int gesture = argmax(output->data.int8, 4);
```

Key Differences:

DATASETS AND BENCHMARKS

STANDARD EVALUATION TOOLS

Popular TinyML Datasets

| Dataset | Task | Samples | Size | Input Type |
|-------------------|----------------------|-----------------------|---------|-----------------|
| Visual Wake Words | Object detection | 115K images | ~5 GB | 96×96 RGB |
| Speech Commands | Keyword spotting | 105K audio | ~2 GB | 1-sec clips |
| Google Commands | Voice recognition | 65K commands | ~1.4 GB | Audio |
| UCI HAR | Activity recognition | 10K sequences | 25 MB | IMU |
| Anomaly Detection | Industrial | 1K normal + anomalies | 100 MB | Sensor |
| Person Detection | Vision | 5K images | 500 MB | 96×96 grayscale |

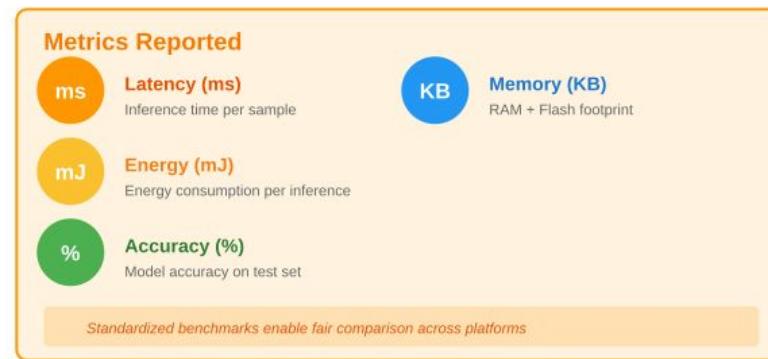
MLPerf Tiny Benchmark

Standard benchmarks for TinyML performance:

| Task | Model | Metric | Result |
|----------------------|-------------|----------|--------|
| Keyword Spotting | DS-CNN | Accuracy | 90.5% |
| Visual Wake Words | MobileNet | mAP | 85.3% |
| Anomaly Detection | Autoencoder | AUC | 0.92 |
| Image Classification | ResNet-8 | Accuracy | 88.7% |

Platforms tested:

Arduino Nano 33 • STM32F746 • ESP32 • NXP i.MX RT



OPTIMIZATION TECHNIQUES

ADVANCED COMPRESSION METHODS

Weight Clustering

Concept: Group similar weights into clusters

```
Original weights:  
[0.23, 0.24, 0.87, 0.89, 1.02, 0.98]  
  
Clustered (2 clusters):  
[0.24, 0.24, 0.88, 0.88, 1.00, 1.00]
```

Benefits:

- ✓ Reduces unique values
- ✓ Improves compression
- ✓ Minimal accuracy loss

Compression: 2-4x

Typical accuracy drop:
< 0.5%

Mixed Precision

Concept: Different layers use different precision

```
Layer 1 (sensitive):      INT16  
Layers 2-5:                INT8  
Layer 6 (sensitive):      INT16  
Average: ~9.5 bits per weight
```

Benefits:

- ✓ Balance accuracy vs size
- ✓ Preserve critical layers
- ✓ Optimize memory usage

Sparse Matrix Compression

Structured Pruning Example:

```
Dense matrix (4x4):  
[0.2 0.1 0.8 0.3]  
[0.0 0.9 0.0 0.4]  
[0.7 0.0 0.0 0.6]  
[0.0 0.5 0.2 0.0]  
  
Compressed (CSR format):  
Values: [0.2, 0.1, 0.8, 0.3, 0.9, 0.4, ...]  
Indices: [0, 1, 2, 3, 1, 3, 0, 3, ...]  
Pointers: [0, 4, 6, 8, 10]
```

Size reduction: 40% → 62.5% savings

Neural Architecture Search

Automated model design for constraints:

```
Objective:      Maximize accuracy  
Subject to:  
• Latency < 10 ms  
• Model size < 100 KB  
• RAM < 20 KB
```

Approaches:

- Reinforcement learning
- Evolutionary algorithms
- Gradient-based NAS

Result:

Custom architectures
optimized for hardware

PROFILING AND POWER MEASUREMENT

INTERACTIVE PERFORMANCE ANALYSIS

Profiling Tools



Energy Trace

(Texas Instruments)

Features:

- Real-time power monitoring
- Energy measurement

Resolution:
4 kHz sampling, nA



STM CubeMonitor

(STMicroelectronics)

Features:

- Variable monitoring
- Power analysis

Use cases:
Debug optimization



Edge Impulse Profiler

Features:

- On-device profiling
- Layer-wise breakdown

Metrics:
Time, RAM, Flash

Example Profiling Output

| Layer | Time (ms) | RAM (KB) | Flash (KB) | Energy (μ J) |
|----------|-----------|----------|------------|-------------------|
| Conv1D_1 | 2.3 | 4.2 | 1.8 | 115 |
| MaxPool | 0.5 | 2.1 | 0.1 | 25 |
| Conv1D_2 | 4.1 | 6.8 | 5.2 | 205 |
| Dense | 1.2 | 0.8 | 2.1 | 60 |
| Total | 8.1 | 13.9 | 9.2 | 405 |

Power Consumption Comparison



Energy Optimization Strategy

1. Reduce inference frequency
2. Use interrupt-driven sensing
3. Optimize model size
4. Sleep between operations

ML COMPILATION FRAMEWORKS

OPTIMIZATION AND CODE GENERATION

MLIR

Multi-Level IR

MLIR

Multi-Level Intermediate Representation

Features:

- Reusable compiler infrastructure
- Multiple abstraction levels
- Dialect system for extensibility
- Progressive lowering

Used by: TensorFlow, PyTorch, IREE

TVM

Apache TVM

Tensor Virtual Machine

Features:

- End-to-end compiler stack
- Auto-tuning for hardware
- Cross-platform deployment
- Supports ARM, x86, GPUs

Best for: Edge devices, mobile, IoT

XLA

XLA

Accelerated Linear Algebra.

Features:

- JIT compilation
- Fusion optimization
- Memory optimization
- TensorFlow backend

Focus: Performance

IREE

ML Runtime

IREE

Intermediate Repr. Execution Engine

Features:

- Built on MLIR
- Ahead-of-time compilation
- Multi-backend support
- Efficient deployment

Target: Mobile, embedded, servers



Glow

Graph Lowering Compiler

Features:

- Neural network compiler
- Hardware-agnostic
- Graph optimization
- PyTorch integration

Focus: PyTorch models

Compilation Stack

High-Level Frameworks

TensorFlow, PyTorch, JAX

IR / Compiler Layer

MLIR, XLA, TVM

Hardware Abstraction

LLVM, Codegen

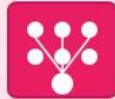
Target Hardware

CPU, GPU, DSP, NPU, MCU

ML HARDWARE ACCELERATORS

SPECIALIZED PROCESSORS FOR EDGE AI

Neural Processing Units (NPUs)



Architecture:

- Dedicated neural network processing
- Parallel MAC (multiply-accumulate) units
- On-chip memory for weights
- Low precision (INT8/INT4) support

Examples:

- Google Edge TPU: 4 TOPS @ 2W
- Apple Neural Engine: 15.8 TOPS
- Qualcomm Hexagon NPU: 15 TOPS

Power Efficiency: 1-5 TOPS/Watt

Digital Signal Processors (DSPs)



Architecture:

- Optimized for signal processing
- Vector instructions (SIMD)
- Hardware loops, zero-overhead
- Fixed-point arithmetic

Examples:

- Cadence Tensilica HiFi: Audio AI
- TI C66x: Industrial, automotive
- ARM Ethos-U: Cortex-M ML extension

Best for: Audio, sensor fusion, control

MCU AI Accelerators



Architecture:

- Cortex-M with ML extensions
- Helium vector extension (M55/M85)
- Integrated MAC accelerators
- Ultra-low power modes

Examples:

- ARM Cortex-M55: 4x ML performance
- NXP i.MX RT: 1 GHz + NPU
- Renesas RA8: Helium + FPU

Power: 10-100 mW during inference

FPGA-based Accelerators



Architecture:

- Reconfigurable logic
- Custom datapath design
- Flexible precision (1-32 bits)
- Pipelined architecture

Examples:

- Lattice sensAI: Ultra-low power
- Intel Cyclone V: Mid-range
- Xilinx Zynq: ARM + FPGA

Best for: Custom algorithms, prototyping

ETHICS AND PRIVACY

LOCAL AI FOR DATA SOVEREIGNTY

Privacy Advantages



Data Stays Local

- ✓ No cloud transmission
- ✓ No third-party access
- ✓ GDPR/HIPAA compliant
- ✓ User control



Security Benefits

- ✓ Reduced attack surface
- ✓ No network vulnerabilities
- ✓ Tamper-evident hardware
- ✓ Encrypted storage

Trade-offs

| Aspect | Cloud AI | Edge AI | TinyML |
|---------------|----------|----------|-----------|
| Privacy | Low | Medium | High |
| Updateability | Easy | Moderate | Hard |
| Debugging | Easy | Moderate | Difficult |

Ethical Considerations

1. Transparency

- How is data processed?
- What decisions are automated?
- Can users understand/audit?
- Explainability requirements

2. Fairness

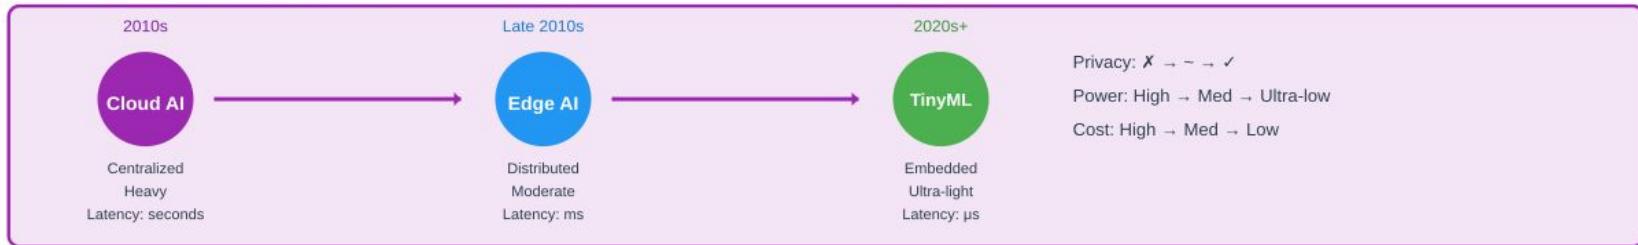
- Bias in training data
- Generalization across groups
- Equitable access to technology
- Inclusive design principles

3. Safety

- Reliability of edge models
- Failure modes and recovery
- Human override mechanisms
- Critical system safeguards

SUMMARY

THE EVOLUTION OF AI



Key Takeaways

1

Efficiency is Everything

- Power constraints drive innovation
- Memory optimization critical
- Compression techniques essential
- 100x reduction while maintaining accuracy
- <1 mW average power achievable

2

Edge AI Enables Real-Time

- Sub-10ms inference possible
- No network dependency
- Deterministic behavior
- Always-on capability
- Offline operation guaranteed

3

Sustainability Through TinyML

- Ultra-low power = battery/solar viable
- Distributed processing reduces load
- Democratizes AI access
- Years of battery life possible
- Reduced carbon footprint

4

Privacy by Design

- Data stays on device
- User sovereignty maintained
- Regulatory compliance built-in
- No cloud data breaches
- GDPR/HIPAA friendly

QUESTIONS?

Dr. Mohammed Alnemari

✉️ malnemar@uci.edu

Thank you for your attention!

Collaboration Welcome

Research partnerships
Industry projects

Key Resources

TinyML Foundation: tinyml.org
Edge Impulse: edgeimpulse.com

Academic Opportunities

Student supervision
Research collaboration